

Agentic AI Gateway — User Manual

Audience: developers who connect applications and coding agents to the gateway, and operators who install, configure, and run it.

Version: September 2026

Contents

1. [What the gateway does](#)
2. [Key concepts](#)
3. [Installing and starting the gateway](#)
4. [First-time setup](#)
5. [Connecting applications](#)
6. [Gateway request options](#)
7. [Grounding answers in your documents \(RAG\)](#)
8. [Letting the gateway run tools \(MCP\)](#)
9. [Using the console](#)
10. [Configuring models and routing](#)
11. [Controlling cost and access](#)
12. [Guardrails](#)
13. [Semantic cache](#)
14. [Monitoring](#)
15. [Errors and what to do about them](#)
16. [Troubleshooting](#)
17. [Reference: settings](#)

1. What the gateway does

The Agentic AI Gateway sits between your AI applications and the large language model (LLM) providers they use. Your applications talk to **one API**; the gateway talks to OpenAI, Anthropic, Google Gemini, and Ollama on their behalf.

Along the way it:

- **keeps requests succeeding** when a provider fails, by retrying and falling back to another provider;
- **chooses the best provider** for each request, by cost, speed, weight, or priority;
- **caches answers** to repeated or near-identical questions;
- **enforces budgets and rate limits** per API key and per team;
- **screens prompts and answers** for personal data and unsafe content;
- **grounds answers in your documents** (RAG) and **runs tools** for the model (MCP);
- **records every request**, with cost, latency, and the route it took.

Applications written for the OpenAI or Anthropic SDKs work unchanged: you only change the base URL and the API key.

2. Key concepts

Term	Meaning
Model name	The name clients ask for, such as <code>gpt-4o</code> . It is a <i>gateway</i> name defined in <code>config/models.yaml</code> , and it can point at any provider.
Deployment	One concrete way to serve a model name: a provider, an upstream model, credentials, and routing settings. A model name can have several deployments.
Virtual key	An API key (<code>sk-...</code>) issued by the gateway to an application. It carries a budget, rate limits, and a list of allowed models. Provider keys are never handed to applications.
Master key	The root key set in <code>MASTER_KEY</code> . It bypasses virtual-key limits. Keep it for administration only.
Team	A group of virtual keys sharing one budget.
Console account	An email and password for a human using the web console. Console accounts and virtual keys are separate; neither works in place of the other.
Fallback	Trying a different deployment after one fails.
Guardrail policy	A named set of rules that block, redact, or flag content.

3. Installing and starting the gateway

What you need

- **Postgres 16** for durable data (keys, logs, users).
- **Redis Stack** (not plain Redis): the cache and RAG need its vector search.
- At least one provider: an API key for OpenAI, Anthropic, or Gemini, or a running Ollama server. The gateway starts with whatever is configured.

Option A: Docker Compose (recommended for a first run)

```
cp .env.example .env          # add your provider keys and secrets
docker compose -f deploy/docker/compose.yaml up -d
```

This starts the gateway on port **8000**, with Postgres, Redis Stack, Prometheus (9090), Grafana (3000), and Ollama (11434).

Option B: from source

```
uv venv --python 3.12
uv pip install -e ".[dev]"
cp .env.example .env
docker compose -f deploy/docker/compose.yaml up -d  # datastores
uv run aigateway migrate                             # create the schema
uv run aigateway serve --port 4000
```

Option C: Kubernetes

A Helm chart is in `deploy/helm/aigateway`. It runs database migrations in an init container before each release. Provide provider keys, `MASTER_KEY`, and `JWT_SECRET` as Kubernetes secrets.

Checking that it is up

URL	Expected
<code>/healthz</code>	<code>200</code> — the process is alive
<code>/readyz</code>	<code>200</code> — the database and Redis are reachable
<code>/ui</code>	The console sign-in page

Port numbers in this manual. Examples use `localhost:4000` (running from source). The Docker image listens on `8000`.

4. First-time setup

4.1 Set the secrets

Before exposing the gateway, set these in `.env` or your secret store:

Setting	Why
<code>MASTER_KEY</code>	Root API key. Anyone holding it can do anything.
<code>JWT_SECRET</code>	Signs console sessions. Use at least 32 random bytes, and the same value on every replica.
<code>BOOTSTRAP_ADMIN_EMAIL</code> / <code>BOOTSTRAP_ADMIN_PASSWORD</code>	Creates the first console administrator, only when no console accounts exist.

If `BOOTSTRAP_ADMIN_PASSWORD` is left at its default of `admin`, the gateway logs a security warning at startup. Change it.

4.2 Create the first administrator from the command line (production)

This keeps the password out of the running process's environment:

```
aigateway create-admin --email admin@example.com --password "$ADMIN_PASSWORD" --role admin
```

4.3 Issue a key for your first application

From the console (**Keys** → **Create key**, see section 9.6), or from the command line:

```
aigateway create-key --name my-app --budget 25 --rpm 60 --model gpt-4o
```

The full key is printed **once**. Only a hash is stored, so copy it now.

4.4 Command-line reference

Command	Purpose
<code>aigateway migrate</code>	Create or upgrade the database schema
<code>aigateway db-stamp</code>	Adopt a database created before migrations were used
<code>aigateway create-admin</code>	Create a console account
<code>aigateway create-key</code>	Mint a virtual key (<code>--budget</code> , <code>--rpm</code> , <code>--model</code> repeatable)
<code>aigateway routes</code>	List the gateway's HTTP routes
<code>aigateway serve</code>	Run the gateway (<code>--port</code> , <code>--reload</code>)

5. Connecting applications

5.1 OpenAI-compatible clients

Endpoints: `/v1/chat/completions` , `/v1/completions` , `/v1/embeddings` , `/v1/models`

Authentication: `Authorization: Bearer <virtual key>`

```
from openai import OpenAI

client = OpenAI(base_url="http://localhost:4000/v1", api_key="sk-your-virtual-key")
reply = client.chat.completions.create(
    model="gpt-4o",
    messages=[{"role": "user", "content": "Hello"}],
)
print(reply.choices[0].message.content)
```

Streaming (`stream=True`), tool calling, JSON mode, and embeddings all work through the same client.

5.2 Anthropic-compatible clients

Endpoints: `/v1/messages` , `/v1/messages/count_tokens`

Authentication: `x-api-key: <virtual key>`

```
from anthropic import Anthropic

client = Anthropic(base_url="http://localhost:4000", api_key="sk-your-virtual-key")
message = client.messages.create(
    model="claude-sonnet-4",
    max_tokens=1024,
    messages=[{"role": "user", "content": "Hello"}],
)
```

Either dialect can reach any model. An Anthropic client can ask for a model served by Gemini, and the gateway translates in both directions.

5.3 Coding agents

Tool	Setting
Claude Code	<code>export ANTHROPIC_BASE_URL=http://localhost:4000</code> and <code>export ANTHROPIC_API_KEY=sk-your-virtual-key</code>
Cursor / Continue	OpenAI base URL <code>http://localhost:4000/v1</code> , API key = your virtual key
opencode	Provider <code>@ai-sdk/openai-compatible</code> with baseURL <code>http://localhost:4000/v1</code>

5.4 curl

```
curl http://localhost:4000/v1/chat/completions \
  -H "Authorization: Bearer sk-your-virtual-key" \
  -H "Content-Type: application/json" \
  -d '{"model": "gpt-4o", "messages": [{"role": "user", "content": "Hello"}]}'
```

Add `"stream": true` and `curl -N` for a streamed reply.

5.5 What the response headers tell you

Header	Meaning
X-Gateway-Provider	Provider that actually served the request
X-Gateway-Deployment	Deployment that served it, such as <code>openai/gpt-4o#2</code>
X-Gateway-Cache	hit or miss
X-Gateway-Cache-Similarity	How close the cached question was (hits only)
X-Gateway-Cost-USD	Estimated cost of this request (0 on a cache hit)
X-Gateway-RAG-Sources	Document chunks used for a RAG request
X-Gateway-Request-Id	Use this to find the request in the console logs

If `X-Gateway-Provider` is not the one you expected, the gateway fell back to another provider. The request log says why.

6. Gateway request options

Any request, in either dialect, can carry extra fields that the gateway uses and removes before calling a provider. Put them at the top level, or under `aigw` if your SDK rejects unknown fields (in the OpenAI Python SDK, use `extra_body={"aigw": {...}}`).

Field	Effect
<code>no_cache</code>	Skip the cache for this request
<code>cache_ttl</code>	Keep this answer in the cache for the given number of seconds
<code>fallbacks</code>	Ordered list of other models to try if this one fails
<code>routing_strategy</code>	<code>priority</code> , <code>least-cost</code> , <code>lowest-latency</code> , <code>weighted</code> , or <code>conditional</code>
<code>guardrail_policy</code>	Use a named guardrail policy
<code>tags</code>	Labels recorded in the log; also steer <code>conditional</code> routing
<code>rag</code>	Ground the answer in a document collection (section 7)
<code>mcp</code>	Let the gateway run tools for the model (section 8)

```
client.chat.completions.create(  
    model="gpt-4o",  
    messages=[{"role": "user", "content": "Summarise this week"}],  
    extra_body={"aigw": {"routing_strategy": "least-cost", "fallbacks": ["claude-sonnet-4"]}},  
)
```

7. Grounding answers in your documents (RAG)

7.1 Create a collection and add documents

In the console, open **RAG** → **Create collection**, then upload documents (see section 9.13). Or use the API:

```
# Create a collection
curl -X POST http://localhost:4000/v1/rag/collections \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"name":"handbook","description":"Staff handbook"}'

# Add a document (JSON body, or a multipart file upload)
curl -X POST http://localhost:4000/v1/rag/collections/<id>/documents \
  -H "Authorization: Bearer $KEY" -H "Content-Type: application/json" \
  -d '{"name":"refunds.txt","content":"Refunds are issued within 14 days..."}'
```

Each document is split into chunks, embedded, and stored. Text documents are supported today.

7.2 Check what will be retrieved

Use the console's **retrieval test** (Figure 9.14), or `POST /v1/rag/search`, to see which chunks match a question and how well they score. If results are poor, the chunk inspector usually shows why: a chunk boundary split a table, or a document ingested as one block.

7.3 Ask a grounded question

```
reply = client.chat.completions.create(
    model="gpt-4o",
    messages=[{"role": "user", "content": "What is our refund policy?"}],
    extra_body={"aigw": {"rag": {"collection_id": "<id>", "top_k": 4}}},
)
reply.aigw["sources"]    # the chunks the answer was based on
```

Options inside `rag`: `top_k`, `min_score`, `mode` (`system` or `user`), `max_context_tokens`, and metadata `filters`. Retrieved text passes through the input guardrails, so a document cannot slip content past your policies.

8. Letting the gateway run tools (MCP)

The gateway can connect to Model Context Protocol (MCP) servers, offer their tools to the model, run the tools the model calls, and feed back the results until the model answers.

8.1 Register a server

In the console, open **MCP** → **Add server** (section 9.14), and choose HTTP (a URL) or stdio (a command run on the gateway host). Registering triggers tool discovery; if it fails, the console shows why. Registering or changing servers needs the master key or a console admin, because a stdio server runs a command on the host.

8.2 Use the tools in a chat request

```
reply = client.chat.completions.create(
    model="gpt-4o",
    messages=[{"role": "user", "content": "What is 2 + 3?"}],
    extra_body={"aigw": {"mcp": {"servers": ["<server id>"], "max_iterations": 8}}},
)
reply.aigw["stop_reason"]          # completed | max_iterations | client_tool_call
reply.aigw["tool_calls_executed"]  # how many tools the gateway ran
```

- Leave out `servers` to offer every healthy server's tools.
- If the model calls a tool your application supplied itself, the gateway stops and returns that call to you (`client_tool_call`).
- Budgets, guardrails, and logging apply once per request; cost is summed across every model call in the loop.

You can also call a tool directly with `POST /v1/mcp/tools/call`. Direct tool calls are not yet logged, costed, or guardrailed.

9. Using the console

Open `http://<gateway-host>/ui` and sign in with a console account. This chapter walks through every screen.

9.1 Roles

Role	Can see	Can change
viewer	Dashboard, Teams, Models, Logs, Usage, Guardrails, Cache, RAG, MCP	Nothing
admin	Everything, plus Keys, Settings, Playground	Everything

The server checks the role on every request; the browser's menus are only a convenience.

9.2 Sessions


- A sign-in lasts one hour and renews itself silently for up to seven days.
- **Sign out** revokes your session on the server immediately.
- Changing your password, or an admin changing your role or deactivating you, signs you out everywhere.
- Change your own password with `POST /admin/api/users/me/change-password` (there is no console screen for it yet). Admins cannot demote or delete their own account, so a gateway is never left without an administrator.

9.3 The console layout

Every screen shares the same frame: the navigation menu on the left (viewers see only the pages they can read), the page title and its main action at the top, and the theme switch and your account menu in the top-right corner. Most pages that report numbers have a **Time window** selector in the top-right.

The screenshots in this chapter come from a demonstration gateway with a few applications, teams, and a morning of mixed traffic.

9.4 Signing in

The image shows a sign-in page for the AI Gateway Console. It features a central white card on a light blue background. The card has the heading "Welcome back" and the instruction "Sign in to your AI Gateway Console." Below this are two input fields: "Email *" with the value "admin@example.com" and "Password *". A blue "Sign in" button is at the bottom of the card.

Welcome back

Sign in to your AI Gateway Console.

Email *
admin@example.com

Password *

Sign in

Figure 9.1: The sign-in page.

Enter your console email and password. Visitors who aren't signed in are sent here from any console address. There is no sign-up link: accounts are created by an administrator (section 9.16) or with `aigateway create-admin`.

9.5 Dashboard




Figure 9.2: The dashboard, showing the last 24 hours.

The landing page answers "is the gateway healthy, and what is it costing?"

- **Top row:** request success rate, p50 and p95 latency, and cache hit ratio, each compared with the previous window.
- **Second row:** total spend, total tokens, **failover count** (requests the gateway rescued by switching deployment), and requests in flight right now.

- **Charts:** requests by status, spend and tokens over time, and requests broken down by model, provider, and key.
- **Provider health:** one chip per deployment with its circuit-breaker state. Green `closed` is healthy. Red `open` means the gateway is skipping that deployment until its cool-down ends. In Figure 9.2, `openai · eval-chat · open` is a deliberately dead deployment, and the failover count shows the gateway routing around it.

A card that says **Not available** means the metric isn't reported for that window. It does not mean zero.

9.6 Keys

Name	Key prefix	Team	Models allowed	Budget / spend	RPM / TPM	Status	Created
legacy-integration	sk-aigw-FWmL...	—	All models	\$0.00 / unlimited	∞ / ∞	Disabled	9/30/2026, 12:52:3.
staging-tests	sk-aigw-3Q8U...	—	All models	\$0.00 / \$5.00	30 / ∞	Active	9/30/2026, 12:52:3.
nightly-summariser	sk-aigw-9fB0k...	—	eval-router, eval-ge...	\$0.00 / \$20.00	∞ / ∞	Active	9/30/2026, 12:52:3.
research-agent	sk-aigw-yg9Q...	Research	All models	\$0.00 / \$150.00	60 / ∞	Active	9/30/2026, 12:52:3.
support-chatbot	sk-aigw-RChYl...	Customer Sup...	All models	\$0.00 / \$50.00	120 / ∞	Active	9/30/2026, 12:52:3.

Figure 9.3: Virtual keys, with their team, allowed models, budget use, rate limits, and status.

Each row shows the key's name, the first characters of its secret (enough to recognise it, not to use it), its team, allowed models, budget and spend, rate limits, and whether it's **Active** or **Disabled**. Each row has action buttons to **edit limits**, **regenerate** the secret, **disable** or **enable** the key, and **delete** it.

To create a key:

1. Select **Create key**.
2. Enter a name. Optionally set a team, budget, rate limits, and allowed models.
3. Select **Create**.




Figure 9.4: Creating a key. Only the name is required.

4. Copy the secret from the confirmation dialog, then select **Done**.




Figure 9.5: The full secret is shown once. Only a hash is stored, so it cannot be shown again.

Use **regenerate** for routine credential changes: it issues a new secret but keeps the budget and history. Use **disable** or **delete** for a leaked key.

9.7 Teams




Figure 9.6: Teams and their shared budgets.

A team gives several keys one shared budget. The page lists each team's budget, spend, and budget period. The usage button shows 30 days of requests, tokens, and cost across the team's keys. Assign a key to a team when you create or edit the key.

9.8 Models and routing

AI Gateway Console

Models & routing
Inspect deployments and the fallback chain used by the gateway. [Reload config](#)

Provider health
anthropic: closed | gemini: closed | ollama: closed | openai: open

Active routing strategy
The backend determines fallback order from deployment configuration; lower priority values are preferred.
priority → health → weighted fallback

Fallback chains

eval-anthropic

1	claude-sonnet-4-5	anthropic	closed	50 ms	Health check
---	-------------------	-----------	--------	-------	------------------------------

eval-multi

1	claude-sonnet-4-5	anthropic	closed	0 ms	Health check
2	gemini-2.5-flash	gemini	closed	0 ms	Health check

eval-gemini

1	gemini-2.5-flash	gemini	closed	1 ms	Health check
---	------------------	--------	--------	------	------------------------------

eval-gemini-embed

1	text-embedding-004	gemini	closed	0 ms	Health check
---	--------------------	--------	--------	------	------------------------------

eval-ollama

1	llama3.2	ollama	closed	1 ms	Health check
---	----------	--------	--------	------	------------------------------

eval-ollama-embed

1	nomic-embed-text	ollama	closed	0 ms	Health check
---	------------------	--------	--------	------	------------------------------

eval-chat

1	gpt-4o-mini	openai	open	0 ms	Health check
2	gpt-4o-mini	openai	closed	1 ms	Health check

eval-router

1	gpt-4o	openai	closed	301 ms	Health check
2	gpt-4o-mini	openai	closed	0 ms	Health check

text-embedding-3-small

1	text-embedding-3-small	openai	closed	0 ms	Health check
---	------------------------	--------	--------	------	------------------------------

Deployment details

Model	Provider	Upstream model	Priority (lower = prefer...)	Weight	Tags	Capabilities	Health
eval-anthropic	anthropic	claude-sonnet-4-5	0	1	native	chat, streaming, tools	closed
eval-multi	anthropic	claude-sonnet-4-5	0	1	native, broken	chat, streaming, tools	closed
eval-gemini	gemini	gemini-2.5-flash	0	1	native	chat, streaming, tool...	closed
eval-gemini-embed	gemini	text-embedding-004	0	1	native, embedding	chat, streaming, emb...	closed
eval-multi	gemini	gemini-2.5-flash	10	1	native, fallback	chat, streaming, tools	closed
eval-ollama	ollama	llama3.2	0	1	native, local	chat, streaming, tool...	closed
eval-ollama-embed	ollama	nomic-embed-text	0	1	native, embedding	chat, streaming, emb...	closed
eval-chat	openai	gpt-4o-mini	0	1	broken, primary	chat, streaming, tool...	open

Rows per page: 10 | 1-10 of 12 | < >

Figure 9.7: Models & routing: provider health, the fallback chain for every model, and deployment details.

This page shows what the gateway will do with a request for each model name:

- **Provider health**, summarised per provider.
- **Fallback chains**: for each model, the deployments in the order they will be tried (1 first), with the upstream model, provider, circuit-breaker state, and recent latency. **Health check** sends a real probe to that deployment, which is the quickest way to confirm a credential works.

- **Deployment details:** every deployment's priority, weight, tags, and capabilities.
- **Reload config** (top-right) re-reads the model catalogue without a restart.

In Figure 9.7, `eval-multi` falls back from an Anthropic deployment to a Gemini one: a cross-provider fallback chain.

9.9 Logs

Time	Request ID	Key	Model	Provider	Status	Latency	Tokens in	Tokens out	Cost
12 seconds ago	req_6597189eef524a...	249313ff-7	eval-anthropic	anthropic	success	54.6 ms	4	6	\$
13 seconds ago	req_5420dddb11ac42...	289e260c-0	eval-chat	openai	success	0.6 ms	12	12	\$
13 seconds ago	req_a614c186b38148...	249313ff-7	eval-chat	openai	success	0.9 ms	6	9	\$
41 seconds ago	req_00f5565f8c8146...	289e260c-0	eval-router	openai	success	1,502.9 ms	7	10	\$
43 seconds ago	req_54310db667b74...	289e260c-0	eval-router	openai	error	208.5 ms	0	0	\$
43 seconds ago	req_937ef9da30e744...	249313ff-7	eval-chat	openai	error	9.9 ms	0	0	\$
43 seconds ago	req_46c8bdf6c32f4b...	289e260c-0	eval-chat	unknown	error	2.3 ms	0	0	\$
43 seconds ago	req_97428603effd44...	289e260c-0	eval-chat	unknown	error	2.5 ms	0	0	\$
43 seconds ago	req_beda7039fc7946...	249313ff-7	eval-chat	openai	success	0.6 ms	15	18	\$
43 seconds ago	req_h8975h7dd4464...	249313ff-7	eval-chat	openai	success	0.6 ms	13	16	\$

Figure 9.8: The request log, newest first.

Filter by time range, virtual key, model, provider, status, cache verdict, minimum latency, or request ID. To find one specific request, paste the `X-Gateway-Request-Id` header from its response into the request ID filter. **Export CSV** downloads every row that matches the filters, not only the visible page.

Select a row to open its details:




Figure 9.9: A request's details: time spent in each pipeline stage, the routing decision, the cache verdict, and guardrail results.

The panel shows:

- **Stage timings:** where the time went (auth, RAG, cache, the provider call, guardrails).
- **Routing decision:** the strategy, the deployment chosen, and why.
- **Cache verdict:** hit or miss, and the similarity score for a hit.
- **Guardrail results:** which rules ran and whether anything was blocked, redacted, or flagged.
- **Attempts:** every retry and fallback. A request that succeeded after two providers failed looks normal in the list; the attempts show that the gateway absorbed an outage.
- **Reveal:** prompt and reply bodies, when body logging is enabled and your role allows it.

9.10 Usage




Figure 9.10: Usage and costs over 30 days, grouped by key, with each key's budget use.

- **Spend and tokens over time** for the selected window.
- **Usage breakdown:** requests, successes, errors, cache hits, tokens, and spend. Use **Group by** to choose key, model, provider, team, or day, and **Export CSV** for finance reports.
- **Key budgets:** each key's spend against its budget. Keys above 80% of their budget are highlighted.

Costs are estimates from the price table (section 10.5).

9.11 Guardrails

AI Gateway Console

Guardrails

Review active policies and recent policy violations.

Policies

default Enabled
block-prompt-injection • redact-pii • input-length • redact-leaked-api-keys

judged Enabled
judge-safety

strict Enabled
block-prompt-injection • redact-pii • strict-input-length • redact-leaked-api-keys

Recent violations

Timestamp	Key	Policy	Rule	Action	Details
9/30/2026, 12:52:32 PM	289e260c-0098-...	default	block-prompt-injection	block	View
9/30/2026, 12:52:32 PM	289e260c-0098-...	default	block-prompt-injection	block	View
9/30/2026, 12:52:32 PM	249313ff-7fda-46...	default	redact-pii	redact	View
9/30/2026, 12:52:32 PM	249313ff-7fda-46...	default	redact-pii	redact	View

Rows per page: 10 ▾ 1-4 of 4 < >
[Previous](#) [Next](#)

Figure 9.11: Guardrail policies and recent violations.

The **Policies** cards list each policy and its rules. **Recent violations** shows each time a rule fired: when, which key, which policy and rule, and the action taken (`block` , `redact` , or `flag`). Select **View** for the details. Figure 9.11 shows two blocked prompt-injection attempts and two prompts where personal data was redacted.

9.12 Cache




Figure 9.12: The semantic cache: effectiveness, stored entries, and invalidation.

- **Top cards:** hit ratio, hits, misses, entries, estimated cost saved, and whether the cache is available.
- **Cache hit ratio over time**, and the configured **semantic threshold**.
- **Entry inspector:** each stored answer with its model, namespace, age, and time to live. Select **Invalidate** to drop one entry.
- **Invalidation:** drop one entry by key, a whole **namespace** (do this after changing a system prompt), or **Flush all cache entries**.

9.13 RAG

The screenshot displays the 'AI Gateway Console' interface. On the left is a sidebar with navigation links: Dashboard, Keys, Teams, Models, Logs, Usage, Guardrails, Cache, RAG (highlighted), MCP, Playground, and Settings. The main content area is titled 'RAG collections' with a subtitle 'Manage indexed knowledge, inspect retrieval results, and test score thresholds.' and a 'Create collection' button. Below this, a dropdown menu shows 'Collections by recency' set to 'Newest first'. A collection named 'Staff handbook' is highlighted, showing '3 documents' and '3 chunks'. To the right, a detailed view of the 'Staff handbook' collection is shown, including its embedding model 'text-embedding-3-small' and dimensions '256 dimensions'. It also displays 'Chunk size 1000 chars' and 'Overlap 150 chars'. Below this, there are tabs for 'Documents', 'Chunks', and 'Retrieval test'. The 'Documents' tab is active, showing an 'Add documents' section with 'Upload file' and 'Paste text' options. The 'Upload file' section includes a 'Chunking strategy' dropdown set to 'Recursive', and input fields for 'Chunk size (characters)' (1000) and 'Overlap (characters)' (150). A button 'Choose .txt or .md file' is present. Below the upload section, a dropdown for 'Documents by recency' is set to 'Newest first'. At the bottom, a table lists documents, with one document 'travel-expenses' shown in detail, including its size (105 B), number of chunks (1), creation and update timestamps, and ingestion status (ready).

Figure 9.13: A RAG collection and its documents.

The left column lists collections. Select one to see its embedding model and chunking settings, with three tabs:

- **Documents:** upload files or paste text, see each document's ingestion status, and delete documents. A document that fails to ingest shows the reason.
- **Chunks:** exactly what was indexed. Check here when retrieval returns something odd.
- **Retrieval test:** run a query and see the scored results before any application depends on the collection.




Figure 9.14: A retrieval test. The best-matching chunk is shown with its similarity score and source document.

9.14 MCP




Figure 9.15: An MCP server with its health and discovered tools.

Select **Add server** to register an HTTP or stdio MCP server. Select a server to see its health, when its tools were last discovered, and each tool's parameters. The buttons at the top of the panel:

- **Check health** probes the server.
- **Rediscover tools** refreshes the tool list.
- **Edit** and **Delete** change or remove the server.

Each tool has a **tool tester**: fill in the parameters and select **Invoke tool** to call it directly.

9.15 Playground




Figure 9.16: The playground: parameters on the left, the conversation in the middle, and the response inspector on the right.

Use the playground to try a model and see exactly how the gateway handled the request. It is admin-only because it spends real money.

- **Parameters (left):** model, system prompt, temperature, top-p, max tokens, stop sequences, and streaming. Under **Gateway controls** you can bypass the cache, override the routing strategy, set a fallback chain, and choose a guardrail policy.
- **Conversation (middle):** type a message and press Enter or **Send**. Replies stream in as they're generated. Use the pencil to edit and resend a message, the arrow to regenerate a reply, and **Clear** to start again.
- **Response inspector (right):** the request ID, latency, time to first token, deployment and provider that answered, cache result, retries, fallbacks, and guardrail results.

Playground requests go through routing, guardrails, caching, and accounting like any other request, and appear in **Logs**.

9.16 Settings

AI Gateway Console

Settings

Review console access and gateway runtime status.

Admin users

Admins can change the gateway; viewers can only read it. [Add user](#)

Email	Role	Status	Created	Actions
finance@example.com	viewer	active	9/30/2026, 12:52:31 PM	Deactivate Delete
ops.viewer@example.com	viewer	active	9/30/2026, 12:52:31 PM	Deactivate Delete
admin@example.com	admin	active	9/30/2026, 12:52:31 PM	You
admin@e2e.test	admin	active	9/30/2026, 12:52:27 PM	Deactivate Delete

Rows per page: 10 1–4 of 4

System information

Version 0.1.0	Environment Gateway runtime	Uptime 0d 0h 0m	Database Connected	Redis Connected
------------------	--------------------------------	--------------------	-----------------------	--------------------

Configured providers

anthropic gemini ollama openai

anthropic	credentials configured	reachable
gemini	credentials configured	reachable
ollama	credentials configured	reachable
openai	credentials configured	unreachable

Optional subsystems

cache: active guardrails: active rag: active mcp: active observability: active

Figure 9.17: Settings: console accounts, system information, provider status, and optional subsystems.

- **Admin users:** select **Add user** to create an account; change a user's role in the **Role** column, or deactivate or delete them. Your own row is marked **You** and can't be demoted or deleted.
- **System information:** version, uptime, and whether the database and Redis are connected.
- **Configured providers:** whether each provider has credentials and is reachable. Secrets are never shown, and they can't be edited from the browser.
- **Optional subsystems:** whether the cache, guardrails, RAG, MCP, and observability are active.

10. Configuring models and routing

10.1 The model catalogue

Models are defined in `config/models.yaml` (path set by `MODELS_CONFIG_PATH`):

```
model_list:
  - model_name: gpt-4o                # what clients ask for
    params:
      provider: openai                # openai | anthropic | gemini | ollama
      model: gpt-4o                   # the provider's model id
      api_key: ${OPENAI_API_KEY}      # expanded from the environment
      priority: 1                     # LOWER number = tried first
      weight: 1
      tags: [premium]

  - model_name: gpt-4o                # same name: a fallback on another provider
    params:
      provider: anthropic
      model: claude-sonnet-4
      api_key: ${ANTHROPIC_API_KEY}
      priority: 2

aliases:
  gpt-4-turbo: gpt-4o                # keep old client names working
```

- Several deployments can share a model name. That is how fallback and load balancing work.
- `${VAR}` and `${VAR:-default}` read from the environment, so secrets stay out of the file.
- **priority is lowest-number-wins.** `priority: 0` is tried before `priority: 10`. Getting this backwards sends traffic to the wrong provider.
- Apply changes with **Models & routing** → **Reload config**, or `POST /admin/api/config/reload`.

10.2 Routing strategies

Strategy	Picks	Good for
priority (default)	Lowest priority number	A clear first choice
least-cost	Cheapest for the request	Background and bulk work
lowest-latency	Fastest recently	Interactive use
weighted	Random, in proportion to <code>weight</code>	A/B tests, gradual rollouts
conditional	Based on the request (length, tools, tags)	Mixed workloads

Set the default with `ROUTING_STRATEGY`; override per request with `routing_strategy`.

10.3 How failures are handled

1. A retryable error (rate limit, timeout, 5xx) is **retried** on the same deployment, up to `MAX_RETRIES` times, with randomised backoff.
2. After that, or for errors like "context too long", the gateway **falls back** to the next deployment, up to `MAX_FALLBACKS` .
3. A deployment that fails `CIRCUIT_BREAKER_THRESHOLD` times in a row is **skipped** for `CIRCUIT_BREAKER_COOLDOWN_SECONDS` , then tested with one request.
4. **Streams fail over only before the first word is sent.** After that, an error ends the stream, because switching providers would mix two different answers.
5. An unknown model name is rejected with `404` . It is never silently served by another model.

10.4 Per-deployment limits

Give a deployment `rpm_limit` or `tpm_limit` (requests or tokens per minute) to match your provider quota. When a deployment is full, traffic spills over to the next one; when all are full, the client gets `429` .

10.5 Prices

Costs are estimated from per-million-token prices in `config/pricing.yaml` (or `pricing` on a deployment). They are close enough to catch a runaway agent, but will not match a provider invoice to the cent.

11. Controlling cost and access

11.1 Limits on a virtual key

Limit	Effect when reached
Budget (USD, optionally per period)	Requests refused with <code>402 budget_exceeded</code>
Requests / tokens per minute	<code>429 rate_limit_exceeded</code>
Requests in flight	Extra concurrent requests get <code>429</code>
Allowed models	Other models get <code>403</code> ; empty list means all
Allowed routes	Restrict a key to, for example, embeddings only
Expiry	The key stops working after the date

11.2 Teams

Put keys in a team to share one budget across them. A request must fit both the key's and the team's budget. Deleting a team keeps its keys working; they stop sharing its budget.

11.3 Good practice

- Give each application its own key, so usage and failures are attributable.
- Keep the master key for administration; never ship it in an application.
- Regenerate keys on a schedule; disable immediately if one leaks.

12. Guardrails

Policies live in `config/guardrails.yaml`. A `default` policy is required; keys can be bound to others, and requests can choose one with `guardrail_policy`.

```
policies:
  default:
    input:
      - type: pii
        action: redact          # block | redact | flag
        entities: [email, phone, credit_card, ssn]
      - type: denylist
        action: block
        terms: [...]
    output:
      - type: pii
        action: redact
```

Rule type	Checks for
<code>pii</code>	Emails, phone numbers, card numbers, national IDs
<code>regex</code>	Any pattern you define
<code>denylist</code>	Forbidden words or phrases
<code>llm_judge</code>	A model scores the text against a policy written in plain language

Action	Effect
<code>block</code>	Refuse with <code>422 guardrail_violation</code>
<code>redact</code>	Replace the matched text and continue
<code>flag</code>	Allow, but record a violation

- Output rules also apply to **streamed** replies. The gateway holds back the last 128 characters (`GUARDRAILS_STREAM_HOLDBACK_CHARS`), so a match split across chunks is still caught.
- An `llm_judge` rule adds a model call to every request under that policy; bind it only to the keys that need it. `on_error` chooses whether a judge failure allows or blocks.
- Watch the violation feed. A rule that constantly fires on normal traffic teaches people to work around it.

13. Semantic cache

The cache answers a request from a stored reply when the question is close enough in meaning to one asked before, not only when the text is identical.

- **Threshold:** `CACHE_SIMILARITY_THRESHOLD` (default 0.95). Lowering it raises the hit rate *and* the risk of answering a question nobody asked. Treat it as a correctness setting.
- **Scope:** entries only match within the same model, key or team, system prompt, temperature range, response format, and tool set.
- **Not cached:** tool calls, and replies generated above `CACHE_MAX_TEMPERATURE` (default 0.3).
- **Skip it:** send `no_cache: true`.
- **After changing a system prompt,** invalidate that namespace from the Cache page; old answers are now wrong in a way no expiry time will catch.

14. Monitoring

Signal	Where
Metrics	Prometheus at <code>/metrics</code> (enable with <code>METRICS_ENABLED</code>)
Dashboards	Grafana, preconfigured in the Docker Compose stack
Traces	OpenTelemetry to <code>OTLP_ENDPOINT</code> (enable with <code>TRACING_ENABLED</code>)
Logs	Structured JSON (<code>LOG_FORMAT=json</code>)
Per-request history	Console Logs page

Prompt and reply bodies are logged only if `LOG_REQUEST_BODIES=true` , and are shown in the console only behind an explicit **reveal**.

Worth alerting on: a deployment's circuit breaker staying `open` ; a falling success rate; spend rising faster than usual; a cache hit ratio suddenly dropping to zero.

15. Errors and what to do about them

Errors come back in the envelope of the dialect you used (OpenAI or Anthropic) with a stable `code`.

HTTP	Code	Meaning	What to do
400	<code>invalid_request</code>	The request is malformed	Fix the request body
400	<code>context_length_exceeded</code>	Too long for every candidate model	Shorten the input, or add a larger model as a fallback
401	<code>authentication_error</code>	Missing, wrong, or disabled key	Check the key and header
402	<code>budget_exceeded</code>	Key or team budget used up	Raise the budget or wait for the next period
403	<code>permission_denied</code>	Model or route not allowed for this key	Update the key's allowlist
404	<code>not_found</code>	Unknown model or resource	Check the model name in <code>/v1/models</code>
422	<code>guardrail_violation</code>	Blocked by a guardrail	Rephrase, or review the policy
429	<code>rate_limit_exceeded</code>	Key limit or deployment capacity reached	Back off and retry
502	<code>all_providers_failed</code>	Every deployment failed	Check provider health on the dashboard
503	<code>no_healthy_deployment</code>	Every deployment is unavailable	Check provider credentials and status
503	<code>rag_unavailable</code>	The RAG store is unreachable	Check Redis

16. Troubleshooting

Symptom	Likely cause and fix
<code>/ui</code> returns 404	Console not built. Run <code>cd ui && npm run build</code> , or use the Docker image.
Cannot sign in on a new deployment	The bootstrap admin is only created when no accounts exist. Use <code>aigateway create-admin</code> .
Everyone signed out after a restart	<code>JWT_SECRET</code> changed. Set it explicitly and identically on all replicas.
Requests are slow and costly	A primary deployment's breaker is <code>open</code> and traffic is on the fallback. Check the dashboard's provider health.
Cache never hits	Check <code>CACHE_ENABLED</code> , that Redis is Redis Stack , and that <code>CACHE_EMBEDDING_DIMENSIONS</code> matches the embedding model.
Cache returns the wrong answer	The threshold is too low, or the system prompt changed: raise the threshold or invalidate the namespace.
A provider shows no credential in Settings	Set its <code>*_API_KEY</code> in the environment and restart. Keys cannot be edited from the browser, by design.
Ollama model not found	Pull it on the Ollama host: <code>ollama pull <model></code> .
MCP server registration fails	The console shows the reason: an unreachable URL, a command not on <code>PATH</code> , or a failed handshake.

17. Reference: settings

The most commonly changed settings. Every setting is an environment variable.

Setting	Default	Purpose
<code>DATABASE_URL</code>	local Postgres	Postgres connection (<code>postgresql+asyncpg://...</code>)
<code>REDIS_URL</code>	<code>redis://localhost:6379/0</code>	Redis Stack connection
<code>MASTER_KEY</code>	—	Root API key
<code>JWT_SECRET</code>	—	Console session signing key (≥ 32 bytes)
<code>OPENAI_API_KEY</code> , <code>ANTHROPIC_API_KEY</code> , <code>GEMINI_API_KEY</code>	—	Provider credentials
<code>OLLAMA_BASE_URL</code>	—	Ollama server
<code>MODELS_CONFIG_PATH</code>	<code>config/models.yaml</code>	Model catalogue
<code>ROUTING_STRATEGY</code>	<code>priority</code>	Default routing strategy
<code>MAX_RETRIES</code>	2	Retries per deployment
<code>MAX_FALLBACKS</code>	3	Deployments tried before giving up
<code>REQUEST_TIMEOUT_SECONDS</code>	120	Whole-request timeout
<code>CIRCUIT_BREAKER_THRESHOLD</code>	5	Failures before a deployment is skipped
<code>CIRCUIT_BREAKER_COOLDOWN_SECONDS</code>	30	How long it is skipped
<code>CACHE_ENABLED</code>	<code>true</code>	Semantic cache on or off
<code>CACHE_SIMILARITY_THRESHOLD</code>	0.95	How close a question must be to hit
<code>CACHE_TTL_SECONDS</code>	3600	Cache entry lifetime
<code>GUARDRAILS_ENABLED</code>	<code>true</code>	Guardrails on or off
<code>LOG_REQUEST_BODIES</code>	<code>false</code>	Log prompts and replies (privacy-sensitive)
<code>METRICS_ENABLED</code>	<code>true</code>	Prometheus metrics
<code>TRACING_ENABLED</code>	<code>false</code>	OpenTelemetry traces
<code>JWT_ACCESS_TTL_SECONDS</code>	3600	Console sign-in length
<code>UI_ENABLED</code>	<code>true</code>	Serve the console

The full list is in [docs/configuration.md](#).